

CMSC 491 Programming Assignment - HTTP

Due Date: Sunday April 7th, 2019 - 11:59pm

Description:

The goal of this assignment is to allow you to demonstrate your knowledge of HTTP requests and HTTP responses. We'll be implementing a simple HTTP client and a simple HTTP server running version of HTTP/1.0. The goal is to get familiar with requests generated by your favorite web browser and responses from real web servers.

HTTPClient

- Your client should be named HTTPClient.java should or HTTPClient.py
- You are going to implement two methods of HTTP: **GET** and **PUT**.
- Your client take command line arguments specifying the needed parameters as follow for each operation.
- **Get**
 - Use the HTTP client to request data from a web server. Your client should be able to send requests and receive responses from any web server on the Internet. There is no need to write a server program to communicate with the client with GET command.
 - Your client should accept a single command-line argument: URL, which is the URL that it will request. There is no need to define GET command in the command line here.
 - The URL will be given in the following format:
http://hostname[:port][/path] (only lowercase letters will be used)
 - hostname - the web server's hostname
 - :port - an optional port, if not present, use port 80.
 - path - the path from the web server's main directory to the requested file, if this is not present, then the path is '/'
 - Example: http://egr.vcu.edu/index.html
 - If any of the arguments are incorrect, exit after printing an error message of the form ERR - arg x, where x is the argument number.
 - Treat any URL as valid as long as it starts with http://
 - Connect to the host given in the URL at the port given in the URL thru TCP connection.
 - Handle exceptions as you wish -- for example, if the host doesn't exist or the port given is not open
 - Submit a valid GET request for path/file given in the supplied URL using HTTP/1.0.
 - Your request must include the "**Host**" field (with the proper entry).

- Your request must include the "**Time**" field (with the proper entry).
 - Your request must also include the "**User-agent**" field with the entry 'VCU-CMSC491'.
 - Your request must also include the "**User-name**" field with *your first name* as the entry.
 - No other request fields are needed.
 - Remember to end with extra CR/LF
- Print the HTTP request sent to the server.
 - Note: Print the actual string sent to the server.
- Parse the HTTP response header received from the server and print the following information:
 - Response code (number)
 - Server type
 - If the response code is in the 200-level:
 - Last modified date
 - Number of bytes in the response data
 - Store the received file (for example: index.html from http://www.vcu.edu/) in the current directory.
 - If the response code is in the 300-level:
 - The URL where the file is located
- Print the HTTP response header.
- In case you successfully received and stored a file, open the stored file using any browser to check it.
- Once you have this part of the client working, you should test it with the following two test cases:
 1. Use it to get a file of your choosing from a "real" web server on the Internet. For example,


```
java HTTPClient http://www.cnn.com/index.html
```
 2. Use it to get a file from your own server program (next). For example, your server is running on 172.18.233.83, port number 10003.


```
java HTTPClient http://172.18.233.83:10003/index.html
```

Note that you can create your own index.html file (doesn't to have a real HTML webpage file) or store any of a downloaded index.html files such as cnn index.html file.

- **PUT**

- Use the HTTP client to put data only to your own web server (described next).
- Your client should accept a three command-line argument: PUT command, URL to where the files should be stored in the server, and the local path/filename of the file to transmit.
 - The URL will be given in the following format:


```
http://hostname[:port]
```

(only lowercase letters will be used)

 - hostname - a CS Unix machine (use either atria or sirius)
 - :port - only ports 10000-11000 are open for use.
 - Example: http://172.18.233.83:10010
 - The path/name of local file has the format: [path/]<filename>

- path - the path of the location of the file to be transmitted at the current machine ,if this is not present, then the path is '/'
 - <filename> is the name of the file that you intend to transmit to the server. If doesn't exist, you
 - Similarly, if any of the arguments are incorrect, exit after printing an error message of the form ERR - arg x, where x is the argument number.
- Connect to the host given in the URL at the port given in the URL thru TCP connection.
 - Handle exceptions as you wish -- for example, if the file doesn't exist, host doesn't exist, or the port given is not open
- Submit a valid PUT request for path/file to be transferred to the webserver given in the supplied URL using HTTP/1.0.
 - Your request must include the "**Host**" field (with the proper entry).
 - Your request must include the "**Time**" field (with the proper entry).
 - Your request must also include the "**User-agent**" field with the entry 'VCU-CMSC491'.
 - Your request must also include the "**User-name**" field with *your first name* as the entry.
 - No other request fields are needed.
 - Remember to end with extra CR/LF
- Print the HTTP PUT request sent to the server.
 - Note: Print the actual string sent to the server. Do not send a request and then re-type the request for output.
- If the server is successful in receiving and storing the file, the server sends back a "200 OK File Created" response. Otherwise, the server sends back a "606 FAILED File NOT Created" response.
- Parse the HTTP response header received from the server and print the following information:
 - Response code (number)
 - Server type
- Print the HTTP response header.

HTTPServer

- Your server should be named HTTPServer.java should or HTTPServer.py
- It must accept a single command-line argument: port, which is the port that it will listen on for incoming HTTP requests.
 - If any of the arguments are incorrect, exit after printing an error message of the form ERR - arg x, where x is the argument number.
 - The only error-checking that needs to be done on the port is to ensure it is a positive integer less than 65536.
 - Remember that only ports 10000-11000 are open for use.
- The server should able to handle both HTTP commands: **GET** and **PUT**.
- For each new HTTP request, print the client's IP address, port, and the request type in the format IP:port:request
 Example: 172.18.233.83:63307:GET

- Print each line in the HTTP request.
- Construct a valid HTTP response including status line, any headers you feel are appropriate, and, of course, the requested file in the response body.
- For **GET**, If the server receives for example the "GET index.html HTTP/1.0" request,
 - Sends out "200 OK" to the client, followed by the file ./index.html.
 - If the requested file doesn't exist, the server sends out "404 Not Found" response to the client.
- For **PUT**, if the server receives the "PUT index.html HTTP/1.0" request, it will save the file as ./index.html.
 - If the received file from client is successfully created, the server sends back a "200 OK File Created" response to the client.
 - Otherwise, the server sends back a "606 FAILED File NOT Created" response.
- The HTTP server program should remain running until the user closes it with Ctrl-C. The program should close the client socket after each request has been received (even if the browser is using HTTP/1.1), but leave the welcoming/listening socket open.
- Once you have your server working, you could test it with the following two test cases:
 1. Use a real browser (e.g., firefox) to GET a file from your server
 - Note that you can download an index.html of any site and store in your server for the test.
 2. Use your HTTPClient to GET/PUT a file

VM Linux Machines:

- Open "terminal" window either on Windows, Mac, or Unix machine.
- On terminal, type "ssh 172.18.233.83 -l <eID>" then your VCU password to authenticate. Note that <eID> is your VCU user login.
- To access another machine with different IP, open a new "terminal" and follow the previous step to connect to the class VM machine (172.18.233.83), and then type "ssh 172.16.49.2 -l <eID>" using the same set of credentials you used before.

Rules:

- The only programming networking classes allowed are the basic socket classes that we've used with the examples. For example, java.net.URL is not allowed and urllib2 (Python) is not allowed.
- Your code should run on the VM Linux machines (e.g., 172.18.233.83). Remember that only ports 10000-11000 are open for use
- You are not permitted to work with anyone else (even students not in the class) - all of the coding and documentation must be your own.
- Your program must compile (if Java/C++) and run on the VM Linux machines.
- You must write neat code and document it well. You will lose points for sloppy programs that contain little or no comments.

Hints:

- Look back in your notes to recall how HTTP requests are formatted and terminated.
- Note that `readLine()` in Java strips off newline characters before returning a `String`.
- Use the `equals()` method in Java to compare two `Strings`.

Testing:

A large part of your program's grade will be determined by how well it handles a set of inputs. You should test your program rigorously before submitting. Because your programs will be run and tested using a script, you must format your output exactly as I have described or you will lose points.

The examples below are just examples. I will test your programs rigorously. In particular, I will test your HTTP Client on a wide range of URLs.

Example 1

```
java HTTPServer
Usage: java HTTPServer port
```

Example 2

In this example, after setting up the server (e.g., on 172.18.233.83), the user opened a web browser (e.g., on 172.16.49.2) to the URL `http://172.18.233.83:10003/my/url`

```
cm5c491> java HTTPServer 10003
172.16.49.2:33083:GET
GET /my/url HTTP/1.1
Host: 172.18.233.83:10003
User-Agent: Mozilla/5.0 (X11; U; SunOS sun4u; en-US; rv:1.7) Gecko/20070606
Accept:
text/xml,application/xml,application/xhtml+xml,text/html;q=0.9,text/plain;q=0.8,i
mage/png,*/*;q=0.5
Accept-Language: en-us,en;q=0.5
Accept-Encoding: gzip,deflate
Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7
Keep-Alive: 300
Connection: keep-alive
```

Example 3

```
java HTTPClient
Usage: java HTTPClient URL or java HTTPClient PUT URL path/<filename>
```

Example 4

```
java HTTPClient http://www.egr.vcu.edu/directory/tamer.nadeem/files/foo.txt
GET /directory/tamer.nadeem/files/foo.txt HTTP/1.0
```

Host: www.egr.vcu.edu
Time: Thu, 21 Feb 2019 20:53:29 GMT
User-agent: VCU-CMSC491
User-name: Tamer

200
Apache/2.2.17 (Unix) PHP/5.3.5 mod_ssl/2.2.17 OpenSSL/0.9.8q
Thu, 19 May 2011 19:23:43 GMT
92

HTTP/1.1 200 OK
Date: Thu, 21 Mar 2013 20:53:29 GMT
Server: Apache/2.2.17 (Unix) PHP/5.3.5 mod_ssl/2.2.17 OpenSSL/0.9.8q
Last-Modified: Thu, 19 May 2011 19:23:43 GMT
ETag: "5c-4a3a5f178cdd0"
Accept-Ranges: bytes
Content-Length: 92
Connection: close
Content-Type: text/plain

Submission Materials:

- Make sure your program compiles and executes on Dept's Linux machines.
- Create a "Readme.txt" file for your program that lists how to compile and execute the program. Include your name and your V# as the first line in the Readme.txt.
- You must name your source programs HTTPClient.java/HTTPServer.java, HTTPClient.py/HTTPServer.py, or HTTPClient.cpp/HTTPServer.cpp.
- Submit all files necessary to compile your program.
- Zip all files of your program files and name them prog_assign_1.
- Submit through Blackboard.

Submitting Assignments using Blackboard

- Instructions from the official Blackboard help pages (with screenshots)

- https://help.blackboard.com/Learn/Student/Assignments/Submit_Assignments

- Step-by-step instructions:

- Log in to the course Blackboard page
- Click the **Programming Assignments** link in the left sidebar
- Select the current assignment
- Under **Assignment Submission**, there is an **Attach File** section.
- Click **Browse My Computer**. Your web browser will open a window in which you can select the file to attach.
- Once you have successfully attached a file, its name should appear in the **Attached files** list.
- You may attach as many files as needed.
- When you've finished adding all of the files that are needed for the assignment, click **Submit** at the bottom of the page.

Important: If you click Save as Draft instead of Submit, your assignment will not be turned in.

- After you've clicked **Submit**, an assignment receipt page will display. Click **OK**.

Important: *If you submit your assignment and later realize you have made a mistake, Blackboard will allow you to re-submit the assignment. The time of your last attempt will be counted as your submission date.*